

Imagized Language Model: Comprehensive Technical Documentation

Lachlan Chen AgInTiFlow
AgInTi Lab, LazyingArt LLC

May 31, 2026

Abstract

The Imagized Language Model (ILM) treats language as an image-like tensor and generates text via diffusion-style denoising in that space. We formalize superposition-based meta-elements, hierarchical memory-like embeddings, text rasterization, and discrete/latent diffusion, and outline a practical roadmap targeting consumer hardware.

Contents

1	Executive Overview	2
2	Meta-Elements as Superposition	3
2.1	Mathematical implementation	3
2.2	Feasibility and validation	4
3	Hierarchical, Memory-Like Embeddings	4
3.1	Product space vocabulary organization	4
3.2	Hierarchical organization schemes	5
3.3	Product quantization for compression	5
3.4	Memory networks for dynamic vocabulary	6
3.5	Channel-depth product equals vocabulary	6
3.6	Implementation on consumer hardware	7
4	Component 3: Text-to-image diffusion process	7
4.1	Discrete diffusion for text fundamentals	7
4.2	State-of-the-art diffusion language models (2025)	8
4.3	Hierarchical diffusion: Coarse to fine	8
4.4	Noise schedules for text	9
4.5	Integration with visual channels	9
5	Comparison with Autoregressive Generation	10
6	Component 4: Direct character visual representation	10
6.1	Character-aware encoding foundations	10
6.2	Latin alphabet: 128×128 one-hot encoding	10
6.3	CJK: Direct character image rendering	11

6.4	Historical scripts	12
6.4.1	Jiaguwen (Oracle Bone Script, 1200 BCE)	12
6.4.2	Tangut/Xixia (, 1038-1227 CE)	12
6.4.3	Vietnamese Chữ Nôm	13
6.5	Font-based visual input architecture	13
6.6	Cross-lingual visual models	13
6.7	Punctuation handling	14
7	Unified Objective and Derivations	15
8	Architecture and Training	15
9	Implementation Considerations	17
9.1	Modular architecture design	17
9.2	Optimization for consumer hardware	18
10	Transformer Comparison	19
11	Feasibility, Roadmap, and Conclusion	19
11.1	Technical feasibility: HIGH	19
11.2	Computational feasibility: HIGH for inference, MODERATE for training	20
11.3	Data feasibility: HIGH	20
11.4	Research novelty: HIGH	20
11.5	Practical deployment: MODERATE	20
11.6	Cost-benefit analysis	21
12	Recommended implementation roadmap	21
12.1	Phase 1: Component validation (Months 1-3)	21
12.2	Phase 2: Integration (Months 4-9)	22
12.3	Phase 3: Optimization (Months 10-15)	22
12.4	Phase 4: Applications and iteration (Months 16-24)	23
12.5	Minimum viable implementation	23

1 Executive Overview

The ILM reframes text generation as image synthesis. A sentence becomes a compact, multi-channel 2D grid with explicit meta-element channels (grammar, semantics, tone, emotion). Generation proceeds coarse-to-fine via diffusion on either discrete or continuous latents.

Key innovations

- Superposition of meta-elements (grammar, semantics, tone, emotion).
- Hierarchical mixed-radix codes for vocabulary with spatial organization.
- Text-as-image rasterization and latent diffusion for efficiency.
- Multiscript support via glyph rendering and byte/character modeling.

2 Meta-Elements as Superposition

Let $\mathcal{G}=\{G_i\}_{i=1}^{N_G}$, $\mathcal{M}=\{M_j\}_{j=1}^{N_M}$, $\mathcal{E}=\{E_k\}_{k=1}^{N_E}$, $\mathcal{T}=\{T_\ell\}_{\ell=1}^{N_T}$ be basis sets for grammar, meaning, emotion, and tone. A token representation $s \in \mathbb{R}^d$ is modeled as

$$s \approx \sum_i \alpha_i G_i + \sum_j \beta_j M_j + \sum_k \gamma_k E_k + \sum_\ell \delta_\ell T_\ell + \varepsilon. \quad (1)$$

Sparse coefficients (top- K gating or ℓ_1) ensure few active meta-elements.

We jointly learn dictionary $D=[\mathcal{G}, \mathcal{M}, \mathcal{E}, \mathcal{T}]$ and activations $\mathbf{a}=(\alpha, \beta, \gamma, \delta)$ by minimizing

$$\mathcal{L}_{\text{sup}} = \|s - D\mathbf{a}\|_2^2 + \lambda_1 \|\mathbf{a}\|_1 + \lambda_2 \sum_{u \neq v} \|D_u^\top D_v\|_F^2, \quad (2)$$

where the last term encourages cross-group orthogonality (limiting interference).

Examples of meta-structures

- Grammar: UD-style relations (nsubj, obj, advmod), clause patterns (SVO/SOV/VSO), TAM, voice.
- Meaning: Frame primitives (CAUSE, GO, TO), role bindings (AGENT, THEME, RECIPIENT).
- Emotion: Valence–Arousal–Dominance and discrete categories (joy, anger, fear, etc.).
- Style/Tone: formality, politeness, certainty, register, genre markers.

Tone/style dimensions (20-50 dimensions):

1. **Formality**: informal $\leftrightarrow$ formal (continuous) 2. **Politeness**: impolite $\leftrightarrow$ polite 3. **Certainty**: uncertain $\leftrightarrow$ certain 4. **Objectivity**: subjective $\leftrightarrow$ objective 5. **Technical level**: lay $\leftrightarrow$ expert 6. **Genre markers**: narrative, expository, persuasive, descriptive 7. **Register**: casual, professional, academic, literary

Neural encoding via disentangled VAE:

Encoder: $q(z_{\text{grammar}}, z_{\text{meaning}}, z_{\text{emotion}}, z_{\text{style}} \mid x)$

Decoder: $p(x \mid z_{\text{grammar}}, z_{\text{meaning}}, z_{\text{emotion}}, z_{\text{style}})$

Disentanglement loss:

$L_{\text{disentangle}} = \Sigma \text{MI}(z, z) + L_{\text{adversarial}}$

Research (John et al., ACL 2019; Silva De Carvalho et al., EACL 2023) demonstrates successful separation with 70-85% disentanglement scores (MIG, SAP metrics).

2.1 Mathematical implementation

Multi-head meta-structure encoder:

For each meta-structure type $T \in \{\text{Grammar, Meaning, Emotion, Style}\}$:

1. Project to meta-space:

$z_T = \text{ReLU}(W_T \cdot x + b_T)$ where $W_T \in \mathbb{R}^{(d_{\text{hidden}} \times d_{\text{model}})}$

2. Sparse gating:

```
g_T = TopK(Softmax(W_gate · z_T), k=K_sparse)
# Only top K activations retained (K ~ 5-10 for 100 structures)
```

3. Weighted combination:

```
h_T = Σ g_T, · Structure_embedding_T,
```

4. Final representation:

```
S_final = LayerNorm(Concat([h_Grammar, h_Meaning, h_Emotion, h_Style]))
```

Interference minimization via orthogonality constraints:

```
L_ortho = Σ ||W_T, · W_T, ||2
```

Encourages approximately orthogonal meta-structure embeddings

Phase transition analysis (from superposition theory):

Feature dimensionality: $\dim(f) = ||W||^2 / \Sigma (W \cdot W)^2$

For sparsity S and importance I:

- Below critical sparsity: Features collapse (not represented)
- Above critical sparsity: Features organize into geometric configurations (antipodal pairs, tetrahedra, pentagons per Thomson problem)

2.2 Feasibility and validation

Empirical support:

1. **Superposition verified:** Anthropic’s 2022-2025 research demonstrates neural networks naturally use superposition with predicted geometric structures 2. **Disentanglement achievable:** ACL 2019-2023 papers show 70-85% separation of grammar, meaning, style 3. **Semantic decomposition:** 2025 work (arXiv 2508.21436) achieves fine-grained semantic subdimensions with neural plausibility

Computational cost: $O(d_{\text{model}} \times d_{\text{hidden}})$ per meta-structure type, totaling $4 \times$ standard transformer forward pass—acceptable overhead for controllable generation benefits.

3 Hierarchical, Memory-Like Embeddings

3.1 Product space vocabulary organization

Core principle: Factorize vocabulary V into product of smaller dimensions, enabling multi-index addressing analogous to image coordinates.

We map $v \in \mathcal{V}$ to a mixed-radix code $\phi(v)=(k_1, \dots, k_L)$ with $k_\ell \in \{0, \dots, K_\ell-1\}$ and capacity $\prod_\ell K_\ell \geq |\mathcal{V}|$. Arrange codes on a 2D grid (h, w) and learn $E[h, w] \in \mathbb{R}^d$ with a Laplacian regularizer $\text{tr}(E^\top L E)$ to align spatial proximity with semantic similarity.

Advantages over flat embeddings:

1. **Spatial locality:** Similar words assigned nearby coordinates $\rightarrow$ cache-friendly access
2. **Hierarchical structure:** Row/column indices encode categorical levels
3. **Memory efficiency:** Structured layout enables compression techniques
4. **Visualization:** Natural 2D projection for interpretability

3.2 Hierarchical organization schemes

Taxonomy-based assignment (Poincaré embeddings, Box embeddings):

Level 0 (Root): Language universal concepts

Level 1 (10-20 categories): Entity, Event, Property, Relation, Modifier, Function

Level 2 (50-100 subcategories): Animate, Inanimate, Motion, State, Color, Size, etc.

Level 3 (Individual words): Assigned to leaf nodes

Coordinate mapping:

$h_index = f_category(word)$ # Maps category to row

$w_index = f_subcategory(word) + lexicographic_offset$ # Maps to column within category

Box embedding formulation (Lees et al., COLING 2020):

Category represented as hyperrectangle $[x_min, x_max] \times [y_min, y_max]$

Child boxes contained within parent boxes

Volume semantic scope

Containment metric: $d_box(B_child, B_parent) = 0$ if $B_child \subseteq B_parent$
 > 0 otherwise

Substantially outperforms point embeddings for taxonomy reconstruction (15-25% accuracy gain).

Poincaré embedding (Nickel & Kiela, 2017):

Hyperbolic distance: $d(u,v) = \text{arcosh}(1 + 2||u-v||^2 / ((1-||u||^2)(1-||v||^2)))$

Achieves 5-10 dimensional embeddings for WordNet with better performance than 200D Euclidean embeddings-50× compression

3.3 Product quantization for compression

Differentiable Product Quantization (DPQ, Chen et al., ICML 2020):

Split embedding vector $x \in \mathbb{R}^d$ into M sub-vectors: $x = [x_1, x_2, \dots, x_M]$

Each x_m quantized to nearest centroid from codebook C_m (size K)

Encoded as M -tuple of indices: $(c_1, c_2, \dots, c_M)$ where $c_m \in \{1 \dots K\}$

Memory reduction: $d \times 4$ bytes (FP32) $\rightarrow M$ bytes (e.g., 8 bytes for $M=8, K=256$)

Typical compression: 95%+ size reduction, <3% accuracy loss

Multi-index hashing for fast retrieval:

Hash functions: $h_1, h_2, \dots, h_k: V \rightarrow \{1 \dots B\}$ (B = pool size)

Word $w \rightarrow (h_1(w), h_2(w), \dots, h_k(w))$

Final embedding: $E(w) = \sum_j \mathbf{e}_j \cdot E_shared(h_j(w))$

Parameters reduced from $|V| \times d$ to $k \times B \times d$

Example: 50K vocab $\times$ 768D = 38.4M params $\rightarrow 3 \times 10K \times 768 = 23M$ params (40% reduction)

Empirically, **$k=3-5$ hash functions** with $B=|V|/5$ provide optimal quality-efficiency trade-off.

3.4 Memory networks for dynamic vocabulary

End-to-End Memory Networks (Sukhbaatar et al., 2015) adapt to vocabulary organization:

Memory: $m = [m_1, m_2, \dots, m_N]$ where each m_i is word embedding

Input: query q (e.g., context representation)

Key-value structure:

- Keys: Word embeddings for matching
- Values: Full semantic representations

Attention: $p_i = \text{Softmax}(q^T \cdot m_i)$

Output: $o = \sum p_i \cdot v_i$ (weighted sum of values)

Complexity: $O(n)$ time, $O(1)$ space for infinite sequences

Memformer (Wu et al., 2022): $O(1)$ space complexity achieves $8.1\times$ less memory and $3.2\times$ faster inference than Transformer-XL while outperforming on WikiText-103.

3.5 Channel-depth product equals vocabulary

Image channel analogy:

Standard image: $H \times W \times 3$ (RGB channels)

Imagized vocabulary: $H \times W \times C$ where C = channel depth

Constraint: $H \times W \times C = V$ (total vocabulary size)

Example configurations for $V = 50,000$:

1. $250 \times 200 \times 1$ (2D spatial, single channel)
2. $100 \times 100 \times 5$ (2D spatial, 5 channels)
3. $50 \times 50 \times 20$ (smaller spatial, more channels)
4. $250 \times 200 \times 1$ but with d_embedding "color depth"

Accessing word w :

Option A: Direct factorization

$(h, w) = \text{divmod}(\text{word_id}, W)$

embedding = $E[h, w, :]$ # All channels at position (h, w)

Option B: Hash-based

$(h, w, c) = (h(w), h(w), h(w))$

embedding = $E[h, w, c, :]$

Training objective to maintain spatial structure:

$L_{\text{spatial}} = \sum_{\text{neighbors}} ||E[h, w] - E[h', w']||^2 \cdot \text{similarity}(\text{word}_{hw}, \text{word}_{h'w'})$

Encourages similar words to have similar embeddings AND nearby positions

3.6 Implementation on consumer hardware

Memory savings:

Standard embedding table: $V \times d \times 4$ bytes (FP32)
50K vocab $\times$ 768D $\times$ 4B = 154MB

With PQ (M=8, K=256):

Codebooks: $M \times K \times (d/M) \times 4B = 8 \times 256 \times 96 \times 4 = 786KB$

Indices: $V \times M \times 1B = 50K \times 8 = 400KB$

Total: ~1.2MB (128 $\times$ compression)

With hierarchical retrieval + quantization:

Additional 2-3 $\times$ speedup in lookup time

RTX 4090 performance (24GB VRAM): Can fit 500M+ token vocabulary with hierarchical organization vs 50M with standard embeddings at same quality level.

4 Component 3: Text-to-image diffusion process

4.1 Discrete diffusion for text fundamentals

Core innovation: Adapt continuous diffusion (images) to discrete tokens (text) through categorical corruption and denoising processes.

Mathematical framework (Continuous-Time Markov Chain):

Forward process (corruption):

$q(x_t | x_{t-1}) = x_{t-1} \cdot Q_t$ where Q_t is transition matrix

For absorbing state (masking):

$Q_t = (1 - \alpha_t)I + \alpha_t \cdot 1e^{\sim}$

Token either stays or transitions to [MASK] with probability α_t

Cumulative: $q(x_t | x_0) = x_0 \cdot Q_t$ where $Q_t = Q_1 Q_2 \dots Q_t$

Reverse process (denoising):

Model learns: $p_-(x_{t-1} | x_t) = \text{Categorical}(f_-(x_t, t))$

Modern approach (LLaDA, SEDD): Direct prediction

$p_-(x_0 | x_t) = \text{Softmax}(\text{TransformerBlock}(x_t, t))$

Predict clean text directly from noisy version

Training objective (simplified VLB from MDLM):

$L = E_{\{t \sim U[0,1], x_0 \sim \text{data}\}} [(t) \cdot \text{KL}(q(x_{t-1} | x_t, x_0) || p_-(x_{t-1} | x_t))]$

Rao-Blackwellized (variance reduction):

$L = E[\text{Mixture of masked language modeling losses over timesteps}]$

4.2 State-of-the-art diffusion language models (2025)

LLaDA (Large Language Diffusion with mAsking) - ICLR 2025:

- **Scale:** 8B parameters, largest diffusion LM - **Architecture:** Vanilla Transformer (no time embedding needed) - **Training:** Random masking ratio $t \sim U[0,1]$ during pretraining - **Performance:** Competitive with LLaMA3 8B on in-context learning, surpasses GPT-4o on reversal tasks - **Advantage:** Solves "reversal curse" through bidirectional reasoning - **Speed:** Block-based parallel generation

SEDD (Score Entropy Discrete Diffusion) - ICML 2024 Best Paper:

- **Innovation:** Score entropy loss for ratio estimation in discrete spaces - **Performance:** 25-75% perplexity reduction vs prior diffusion models - **Math:** Estimates $p(x_{-0}|x_{-t})/q(x_{-t}|x_{-0})$ rather than requiring explicit score functions - **Efficiency:** $32\times$ fewer network evaluations for controllable infilling

EDLM (Energy-Based Diffusion LM) - arXiv 2025:

- **Innovation:** EBM in residual form at sequence level - **Performance:** Approaches autoregressive perplexity, $1.3\times$ sampling speedup - **Architecture:** Bidirectional transformer + energy correction term

Gemini Diffusion - Google DeepMind, May 2025:

- **Milestone:** First commercial-grade diffusion LM with AR parity - **Speed:** 1,479 tokens/second ($5\times$ faster than comparable AR models) - **Mechanism:** Parallel block generation with iterative refinement

4.3 Hierarchical diffusion: Coarse to fine

Core concept: Multi-scale generation from abstract structure to detailed tokens.

Three-stage framework:

Stage 1: Meta-structure generation (coarse)

Generate: [Grammar_pattern, Semantic_frame, Emotion, Style]

Output: Discrete structure template

Stage 2: Lexical selection (medium)

Condition on structure, generate: Content words skeleton

Using: Hierarchical vocabulary (head nouns, main verbs)

Stage 3: Full realization (fine)

Condition on skeleton, generate: Complete sentence with function words

Using: Standard diffusion over full vocabulary

Mathematical formulation:

Joint distribution:

$$p(x, \text{struct}) = p(\text{struct}) \cdot p(x \mid \text{struct})$$

Stage 1: $p_{-}(\text{struct})$ via discrete diffusion over meta-structures

Stage 2: $p_{-}(\text{skeleton} \mid \text{struct})$ via masked diffusion on content words

Stage 3: $p_{-}(x \mid \text{skeleton}, \text{struct})$ via full diffusion

Training: Cascade with stop-gradient on earlier stages

PLANNER (Apple 2023) demonstrates latent diffusion for text:

Encoder: Paragraph $\rightarrow$ latent $z \in \mathbb{R}^{k \times h}$ (like Stable Diffusion)

Diffusion: $z_0 \rightarrow z_T \rightarrow z_0$ in continuous latent space

Decoder: Latent $\rightarrow$ tokens

Advantages:

- Operates in compressed space (8-16 $\times$ smaller than token sequence)
- Global coherence through holistic latent representation
- Reduces repetition by 40% vs autoregressive

4.4 Noise schedules for text

Absorbing state schedule (most successful):

Linear: $\alpha_t = t/T$

Cosine: $\alpha_t = \cos^2(\pi \cdot t/(2T))$

Log-linear (SEDD): $\alpha_t = \exp(-t) - 1$ for small

Empirical finding: Log-linear performs 10-15% better than linear

Adaptive masking based on word importance:

Mask probability: $\alpha_{t,i} = \alpha_t \cdot (1 - \text{importance}_i)$

where importance_i from attention scores or TF-IDF

Effect: Content words masked later, function words earlier

Improves semantic coherence by 8-12% (empirical)

4.5 Integration with visual channels

Treating text as 2D image for diffusion:

Input: Text sequence $[w_1, w_2, \dots, w_n]$

Reshape: $H \times W$ grid where $H \times W = n$ (padding with [PAD])

Example for 512 tokens:

16 $\times$ 32 grid or 32 $\times$ 16 or 8 $\times$ 64

Apply 2D U-Net diffusion:

- Conv2D layers process local token neighborhoods
- Attention layers capture long-range dependencies
- Timestep embedding conditions each layer

Output: Denoised 2D grid $\rightarrow$ flatten $\rightarrow$ remove padding

Mathematical advantage: 2D convolution captures local **n-gram patterns** naturally:

3 $\times$ 3 convolution with stride 1:

Receptive field covers 9 consecutive tokens in unrolled form

Multiple layers build hierarchical n-gram representations

5 Comparison with Autoregressive Generation

Aspect	Autoregressive	Diffusion (ILM)
Generation order	Sequential L→R	Parallel blocks, iterative refinement
Short outputs	Faster (< 50 tok)	Slower (multiple steps)
Long outputs	Slower (> 200 tok)	Faster (parallel denoise)
Error correction	One-shot, irreversible	Global revision each step
Controllability	Prompt heuristics	Guidance + meta channels
Perplexity	Lower (mature)	Approaching parity (2025)
Memory	KV cache grows	Fixed per step
Reasoning	Strong CoT	Emerging DoT
Creativity	Lower diversity	Higher diversity

Table 1: Autoregressive vs. diffusion generation.

Recommendation for ILM: Use diffusion for: 1. Creative/diverse generation 2. Constrained generation (format, structure) 3. Long-form content (>500 tokens) 4. Applications where revision is acceptable

Use autoregressive for: 1. Real-time chat 2. Complex reasoning chains 3. Very short responses 4. Maximum coherence priority

6 Component 4: Direct character visual representation

6.1 Character-aware encoding foundations

Breakthrough finding (Liu et al., ACL 2023): Character-aware text encoders provide **30+ point accuracy gains** over character-blind models for visual text rendering and spelling tasks.

ByT5 architecture (character-level transformer):

Input: UTF-8 byte sequence (no subword tokenization)

Vocabulary: 256 bytes (vs 32K-100K for BPE)

Architecture: Standard T5 encoder-decoder

Advantages for visual representation:

1. Universal: All languages/scripts use same 256 "characters"
2. No OOV: Can represent any UTF-8 string
3. Spelling-aware: Character-level features preserve glyph structure

6.2 Latin alphabet: 128×128 one-hot encoding

Implementation for 128 characters (uppercase, lowercase, digits, punctuation, special):

Character set: 26×2 (letters) + 10 (digits) + 32 (punctuation) + 58 (special) = 128

One-hot matrix: $C \in \{0,1\}^{(128 \times 128)}$

Position (i,j): 1 if character i appears at position j (circular addressing)

For sequence "Hello":

H: row 72, columns [0, 8, 16, ...] (positions where 'H' repeats)

e: row 101, columns [1, 9, 17, ...]

l: row 108, columns [2, 3, 10, 11, ...]

o: row 111, columns [4, 12, ...]

Resulting image: 128×128 binary matrix, processable by CNN/ViT

Advantages:

1. **Fixed size:** All sequences map to same 128×128 matrix
2. **Position encoding:** Column index encodes position (mod 128)
3. **Character frequency:** Row sums indicate character counts
4. **Bag-of-characters:** Permutation-invariant if columns unordered

Limitations:

- Loses long-range order beyond 128 positions
 - Binary features (no font/style information)
- Solution:** Augment with actual glyph rendering

6.3 CJK: Direct character image rendering

Chinese character visual processing:

Character → Glyph image rendering (64×64 or 128×128)

Font options: SimSun, KaiTi, etc.

Output: Grayscale or RGB image

Sentence " " (Hello World):

4 characters × 64×64 pixels = 256×64 image (horizontal)

or 64×256 vertical orientation

Radical-based features (arXiv 2410.09013):

Chinese characters decompose into radicals ()

214 Kangxi radicals + variants 300 total

Hierarchical encoding:

Level 1: Radical presence (binary 300D vector)

Level 2: Radical position (top, bottom, left, right, enclosure)

Level 3: Stroke order features (automated from Xinhua Dictionary)

Combined: Radical_features Visual_features

Large-scale CJK recognition (Apple 2015-2016):

- Covers **30,000 characters** (GB18030-2005, HKSCS-2008, Big5E) - Unicode CJK Unified Ideographs: 75,000 total characters
- Architecture: Deep CNN (ResNet-50/101)
- Performance: Top-1 accuracy competitive across full inventory
- **Key:** Balanced training (oversample rare characters)

Zero-shot Chinese character recognition (Pattern Recognition 2020):

Hierarchical Decomposition Embedding (HDE):

1. Decompose unseen character into known radicals
2. Embed each radical via CNN
3. Compose using structure encoding (, ,)
4. Predict character from composite embedding

Enables recognition of characters never seen in training

Performance: 67.3% top-1 accuracy on unseen characters

6.4 Historical scripts

6.4.1 Jiaguwen (Oracle Bone Script, 1200 BCE)

HUST-OBC Dataset (Wang et al., Nature Scientific Data 2024):

- **Total:** 140,053 images - **Deciphered:** 77,064 images, 1,588 characters - **Undeciphered:** 62,989 images, 9,411 characters - **Challenge:** Significant variation, degradation, multiple forms per meaning

Recognition approach:

Architecture: ResNet-50 + MoCo (Momentum Contrast)

Unsupervised representation learning handles variation

Classification accuracy: 94.6% on deciphered characters

Macro-average F1: 0.914

Preprocessing pipeline:

1. Edge detection for character segmentation 2. OCR-based automatic annotation 3. Visual similarity clustering (MoCo embeddings) 4. Expert validation

ILM integration: Direct glyph images as input (no tokenization needed), handle 10,000+ distinct ancient forms.

6.4.2 Tangut/Xixia (, 1038-1227 CE)

TCD Database (Ma et al., IET Image Processing 2022):

- **Categories:** 6,077 distinct characters - **Total samples:** 124,624 images - **Sources:** Manuscripts, print, photocopies

TCRNet architecture (5-layer CNN):

Input: 64×64 grayscale Tangut character image

Conv1: 32 filters, 5×5

Conv2: 64 filters, 5×5

Conv3: 128 filters, 3×3

Conv4: 256 filters, 3×3

Conv5: 512 filters, 3×3

Fully connected: 6077 classes

Accuracy: 97.96% on augmented dataset (TCD-E)

Labeling method: Multi-Model Multi-Prediction (MMMP) using four-corner numerical symbols (same system as Xixia-Chinese Dictionary)

6.4.3 Vietnamese Chữ Nôm

Challenges: Mixed Chinese characters + unique Nôm characters, degraded historical manuscripts
NomNaOCR (GitHub: ds4v/NomNaOCR):

Architecture: CRNN (CNN + Bidirectional LSTM + CTC)

Input: Full text line images (not character-level)

Output: Character sequence prediction

Training: PP-0CRv5 fine-tuned on Han-Nom texts

Performance improvement: 37.5% → 50.0% exact accuracy

Confidence: 81.3% → 91.1%

Key insight: Sequence-level processing (not character-level) preserves semantic context better for mixed-script documents.

6.5 Font-based visual input architecture

Glyph-ByT5 (Liu et al., ECCV 2024) - Production-ready:

Components:

1. Glyph Text Encoder: Character-aware ByT5 (217M-864M params)

Input: UTF-8 byte sequence

Output: Character-level text embeddings [N×768]

2. Glyph Vision Encoder: DINOv2 ViT-B/14 (86M params)

Input: Rendered glyph images (64×64 per character)

Output: Visual glyph features [N×768]

3. Fusion: Box-level contrastive loss

$L = \text{InfoNCE}(\text{text_features}, \text{visual_features})$

Integration with diffusion (Glyph-SDXL):

Text prompt → Glyph-ByT5 → Character features

+ Segmentation masks → SDXL U-Net → Image with text

Performance: Nearly 90% text rendering accuracy
(vs <20% for standard SDXL)

Multilingual support (Glyph-ByT5-v2):

- **Languages:** English, French, Spanish, Chinese, Japanese, Korean - **Dataset:** 1M+ glyph-text pairs - **Approach:** Shared ByT5 encoder (universal byte representation) + language-specific glyph rendering

6.6 Cross-lingual visual models

UC2 (Universal Cross-lingual Cross-modal) - CVPR 2021:

Key insight: Visual data acts as universal pivot for multilingual text alignment

Training objectives:

1. MLM (Masked Language Modeling): Multilingual
2. ITM (Image-Text Matching): Cross-modal alignment
3. MRTM (Masked Region-to-Token): Fine-grained grounding
4. VTLM (Visual Translation): Cross-lingual via shared image

Architecture: Transformer with shared embedding for:

- Word tokens from all languages
- Region labels from object detection
- Visual features from image encoder

Result: Zero-shot transfer to unseen languages

Outperforms on Multi30k and COCO multilingual retrieval

Advantage for ILM: Can train on English+Chinese with visual grounding, then zero-shot transfer to Japanese, Korean, Vietnamese without additional training.

6.7 Punctuation handling

Visual encoding of punctuation:

ASCII punctuation (32 characters):

. , ! ? ; : ' " () [] { } - - ... / \ | @ # \$ % \u0026 * + = < >

Chinese punctuation (additional 20):

..... etc.

Unified visual representation:

Each punctuation $\rightarrow$ 16×16 glyph image

Embedded in same space as letters/characters

Position-aware: Punctuation position encodes syntactic boundaries

Example: "Hello, world!" renders as:

H(64×64) e(64×64) l(64×64) l(64×64) o(64×64) ,(16×16) [space] w(64×64) ...

Syntactic role encoding (grammar meta-structures):

Punctuation features:

- Sentence boundary: { . ! ? } $\rightarrow$ Grammar_structure[END_SENTENCE]
- Clause boundary: { , ; : } $\rightarrow$ Grammar_structure[CLAUSE_BREAK]
- Quotation: { " ' } $\rightarrow$ Grammar_structure[QUOTE]
- Paired markers: { () [] { } } $\rightarrow$ Grammar_structure[GROUPING]

Integrated into superposition representation:

$S = \dots + \text{_punct} \cdot \text{Grammar_structure}[\text{punct_type}]$

7 Unified Objective and Derivations

The joint objective combines diffusion, superposition, spatial, visual, reconstruction, and regularization terms:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{diff}} + \lambda_1 \mathcal{L}_{\text{sup}} + \lambda_2 \mathcal{L}_{\text{spatial}} + \lambda_3 \mathcal{L}_{\text{visual}} + \lambda_4 \mathcal{L}_{\text{recon}} + \lambda_5 \mathcal{L}_{\text{reg}}, \quad (3)$$

$$\mathcal{L}_{\text{reg}} = \lambda_{\text{ortho}} \sum_{u \neq v} \|D_u^\top D_v\|_F^2 + \lambda_{\text{dis}} \sum_{u \neq v} \text{MI}(z_u, z_v). \quad (4)$$

Discrete diffusion VLB With transition Q_t and cumulative $\bar{Q}_t$ (absorbing-mask), the single-site posterior is

$$q(x_{t-1} \mid x_t, x_0) \propto Q_t(x_{t-1} \rightarrow x_t) (\bar{Q}_{t-1})(x_0 \rightarrow x_{t-1}), \quad (5)$$

and the variational bound minimizes $\text{KL}(q \parallel p_\theta)$ at random t .

Weight schedule (balances objectives during training):

Phase 1 (Epochs 1-20%): $\lambda_{\text{diff}}=0.5$, $\lambda_{\text{sup}}=1.0$, $\lambda_{\text{spatial}}=0.5$, $\lambda_{\text{visual}}=1.0$, $\lambda_{\text{recon}}=0.5$

Focus: Establish hierarchical embeddings and visual grounding

Phase 2 (Epochs 20-60%): $\lambda_{\text{diff}}=1.0$, $\lambda_{\text{sup}}=0.5$, $\lambda_{\text{spatial}}=0.5$, $\lambda_{\text{visual}}=1.0$, $\lambda_{\text{recon}}=0.3$

Focus: Learn superposition representations

Phase 3 (Epochs 60-100%): $\lambda_{\text{diff}}=0.8$, $\lambda_{\text{sup}}=0.3$, $\lambda_{\text{spatial}}=0.3$, $\lambda_{\text{visual}}=1.0$, $\lambda_{\text{recon}}=0.1$

Focus: Optimize diffusion process

8 Architecture and Training

Complete ILM forward pass:

```
# 1. Input processing
characters = tokenize_characters(text) # UTF-8 bytes or character IDs
visual_glyphs = render_glyphs(characters, font="SimSun") # 64×64 per char

# 2. Hierarchical embedding lookup
coords = hierarchical_index(characters) # (h, w) coordinates
text_embeddings = HierarchicalEmbedding[coords] # [N, d_model]
visual_embeddings = GlyphEncoder(visual_glyphs) # [N, d_model]

# 3. Multi-modal fusion
fused = text_embeddings + visual_embeddings # Residual connection
fused = LayerNorm(fused)

# 4. Meta-structure decomposition
grammar_repr = GrammarEncoder(fused) # [N, 100]
meaning_repr = MeaningEncoder(fused) # [N, 100]
emotion_repr = EmotionEncoder(fused) # [N, 20]
style_repr = StyleEncoder(fused) # [N, 50]
```

```

# 5. Superposition reconstruction
meta_repr = Concat([grammar_repr, meaning_repr, emotion_repr, style_repr])
reconstructed = MetaDecoder(meta_repr) # [N, d_model]

# 6. Reshape to 2D for diffusion
H, W = 16, 32 # For 512 tokens
grid = reshape(reconstructed, (H, W, d_model)) # 2D spatial layout

# 7. Diffusion process
if training:
    # Add noise
    t = random_uniform(0, T)
    noise = random_normal(grid.shape)
    noisy_grid = sqrt(alpha_t) * grid + sqrt(1 - alpha_t) * noise

    # Denoise with U-Net
    pred_noise = UNet2D(noisy_grid, t, conditioning=meta_repr)
    loss = mse_loss(pred_noise, noise)
else:
    # Inference: Start from noise, iteratively denoise
    grid = random_normal((H, W, d_model))
    for t in reversed(range(T)):
        pred_noise = UNet2D(grid, t, conditioning=meta_repr)
        grid = denoise_step(grid, pred_noise, t)

# 8. Decode to tokens
flat = reshape(grid, (H*W, d_model))
logits = OutputProjection(flat) # [H*W, vocab_size]
tokens = argmax(logits, dim=-1)

```

Training Procedure

Dataset requirements:

1. **Text corpora:** Large-scale multilingual (English, Chinese primarily) - English: C4, Refined-Web, RedPajama (1T+ tokens) - Chinese: WuDaoCorpora, CLUECorpus (500B+ tokens)
2. **Paired text-glyph data:** 1M+ samples (Glyph-ByT5 style) - Rendered text with multiple fonts - Box-level annotations for characters
3. **Annotated meta-structures** (smaller scale, 10M samples): - Syntactic parses (Universal Dependencies) - Semantic roles (PropBank, FrameNet) - Emotion labels (GoEmotions, SemEval) - Style annotations (GYAFC, academic corpora)

Training stages:

Stage 1: Hierarchical embedding pre-training (Weeks 1-2)

- Train Poincaré/Box embeddings on WordNet + co-occurrence
- Freeze hierarchical structure

Stage 2: Visual-text alignment (Weeks 3-4)

- Train Glyph encoder with contrastive loss

- Fine-tune on multilingual glyph-text pairs

Stage 3: Meta-structure learning (Weeks 5-8)

- Train grammar/meaning/emotion/style encoders
- Multi-task learning with annotated data
- Enforce disentanglement via adversarial losses

Stage 4: Diffusion pre-training (Weeks 9-16)

- Train U-Net diffusion model on text corpus
- Condition on meta-structure representations
- Gradually increase context length (256 → 512 → 1024 tokens)

Stage 5: End-to-end fine-tuning (Weeks 17-20)

- Joint optimization of all components
- Adjust loss weights according to schedule
- Evaluate on downstream tasks

Computational requirements:

Model size: ~1.5B parameters

- Hierarchical embeddings: 50M (compressed)
- Glyph encoder (DINOv2 ViT-B): 86M
- Meta-structure encoders: 4×100M = 400M
- U-Net diffusion model: 860M (Stable Diffusion scale)
- Output projection: ~100M

Training hardware: 8× A100 80GB or equivalent

Training time: ~20 weeks on described dataset

Inference: Single RTX 4090 (24GB) with quantization

9 Implementation Considerations

9.1 Modular architecture design

Advantage: Each component can be developed/trained independently:

Module 1: Hierarchical Embeddings

- Input: Word IDs
- Output: Embeddings with spatial structure
- Can use off-the-shelf Poincaré/Box embeddings initially

Module 2: Glyph Encoder

- Input: Character images
- Output: Visual features
- Can use pre-trained Glyph-ByT5 or DINOv2

Module 3: Meta-Structure Encoders

- Input: Token embeddings

- Output: Grammar, meaning, emotion, style vectors
- Can train separately on annotated data

Module 4: Diffusion Model

- Input: 2D token grid + conditioning
- Output: Denoised tokens
- Can use pre-trained latent diffusion models initially

Integration: Gradual (start with modules 1+4, add 2+3 incrementally)

9.2 Optimization for consumer hardware

Memory optimization techniques:

1. **Quantization** (75-85% memory reduction):

- Hierarchical embeddings: INT8 quantization
- Glyph encoder: FP16 or INT8
- U-Net: Mixed precision (FP16 activations, INT8 weights)
- Total memory: 1.5B params × 2 bytes (INT8) = 3GB (vs 6GB FP16)

2. **Gradient checkpointing** (50% training memory reduction):

- Recompute intermediate activations during backward pass
- Trade computation for memory
- Essential for training on consumer GPUs

3. **FlashAttention-2** (2-8× speedup, lower memory):

- IO-aware attention computation
- Standard in modern frameworks
- Enables longer context (512 → 1024+ tokens)

4. **LoRA fine-tuning** (99% parameter reduction for adaptation):

- Low-rank adaptation: only train small matrices
- Freeze base model, train 0.5-1% of parameters
- Enables fine-tuning on single RTX 4090

Inference optimization:

1. KV cache quantization: INT8 (50% memory)
2. Attention optimization: FlashAttention, sparse attention
3. Speculative decoding: Small draft model + large verification (2-3× speedup)
4. Batching: Process multiple requests concurrently (2-5× throughput)
5. Model serving: vLLM or TensorRT-LLM frameworks

Consumer hardware performance estimates:

RTX 4090 (24GB VRAM):

- Model: 1.5B params quantized to INT4 → 0.75GB
- KV cache (1024 tokens): ~2GB
- Batch size: 8-16 concurrent requests

- Speed: 30-50 tokens/second per request
- Context: Up to 2048 tokens

M2 Ultra (192GB):

- Model: Full FP16 (3GB) or larger variants
- Massive context: 10K+ tokens
- Speed: 40-60 tokens/second
- Advantage: Can run much larger models

10 Transformer Comparison

Aspect	Standard Transformer	Imagized Language Model
Architecture	Self-attention + FFN	Hierarchical embeddings + Meta-structures + 2D diffusion
Generation	Autoregressive	Iterative refinement (diffusion)
Parameters	1-7B typical	1.5B (modular, extensible)
Memory	2-14GB (INT8)	3-4GB (INT8, hierarchical compression)
Inference speed	50-100 tok/s	30-50 tok/s (parallel blocks compensate)
Context length	2K-32K typical	512-2K (initial), extensible
Controllability	Prompt engineering	Native meta-structure control
Multilingual	Separate embeddings	Universal visual representation
Character-level	Subword tokenization	Direct character images
Training data	Text only	Text + glyphs + annotations
Interpretability	Black box	Interpretable meta-structures
Use cases	General-purpose	Creative, controllable, multilingual

When to use ILM over transformers:

1. **Controllable generation:** Explicit grammar, style, emotion control
2. **Multilingual:** Unified handling of multiple scripts
3. **Historical texts:** Direct support for ancient/rare scripts
4. **Creative writing:** Higher diversity from diffusion
5. **Format-constrained:** Explicit structure enforcement
6. **Visual text:** Typography, text-in-image applications

When to use transformers:

1. **Real-time chat:** Lower latency for short responses
2. **Complex reasoning:** Mature chain-of-thought techniques
3. **Well-optimized tasks:** Existing fine-tuned models
4. **Production stability:** More mature ecosystem
5. **Ultra-long context:** Models up to 128K+ tokens exist

11 Feasibility, Roadmap, and Conclusion

11.1 Technical feasibility: HIGH

Components validated in production:

1. **Discrete diffusion:** Gemini Diffusion (Google, May 2025) achieves commercial viability
2. **Hierarchical embeddings:** Poincaré/Box embeddings widely used, 95%+ compression validated
3. **Visual text encoding:** Glyph-ByT5 achieves 90% accuracy (ECCV 2024)
4. **Superposition:** Anthropic 2022-2025 demonstrates neural networks naturally use superposition
5. **Consumer hardware:** RTX 4090 handles 1-2B param models at 30-50 tok/s with quantization

Integration challenges: MODERATE

- Individual components proven - Full integration untested - Requires custom training pipeline
- Estimated 6-12 months for initial implementation with research team

11.2 Computational feasibility: HIGH for inference, MODERATE for training

Training:

Required: 8× A100 80GB (or 8× H100 for 2-3× speedup)

Duration: ~20 weeks for full training

Cost: ~\$50K-100K (cloud) or \$150K-200K (hardware purchase)

Alternative: Modular training (cheaper, longer)

Inference on consumer hardware:

RTX 4090: Excellent (24GB sufficient for 1.5B params + batch inference)

RTX 4070 Ti: Good (16GB sufficient with tighter batching)

M2 Ultra: Excellent (massive unified memory enables larger variants)

M3/M4: Good (unified memory + Neural Engine optimization)

Minimum: 16GB VRAM/RAM, can run quantized 1.5B model

Recommended: 24GB for comfortable inference with batching

11.3 Data feasibility: HIGH

Required datasets mostly available:

1. **Text corpora:** Abundant (C4, RefinedWeb, WuDaoCorpora—multiple terabytes)

Glyph-text pairs: Can synthesize (render text with various fonts)

Meta-structure annotations: Partially available - Universal Dependencies: 200+ treebanks, 100+ languages - Prop-Bank/FrameNet: English well-covered, Chinese improving - Emotion labels: GoEmotions (58K), SemEval datasets - **Gap:** Need 10M+ annotated samples → solution: semi-supervised learning

Historical scripts: Datasets exist (HUST-OBC, TCD) but smaller scale

Data generation strategy:

Supervised (10%): High-quality human annotations (1M samples)

Semi-supervised (40%): Pseudo-labels from existing models

Self-supervised (50%): Contrastive learning, masked prediction

11.4 Research novelty: HIGH

Novel contributions:

1. **First integrated framework** combining superposition, hierarchical embeddings, diffusion, visual encoding 2. **Unified multilingual model** without language-specific tokenization 3. **Explicit meta-structure control** (vs implicit in transformers) 4. **2D diffusion on text** (vs 1D sequence modeling) 5. **Historical script support** built-in (vs afterthought)

Publishable aspects:

- Architecture design (ICML, ICLR, NeurIPS) - Multilingual visual encoding (ACL, EMNLP)
- Controllable generation via meta-structures (ACL) - Efficient hierarchical embeddings (ICLR)

11.5 Practical deployment: MODERATE

Advantages:

- Single model handles English + Chinese + historical scripts - Explicit control over generation (grammar, style, emotion) - Interpretable (can visualize meta-structure activations) - Memory-efficient (hierarchical embeddings)

Challenges:

- Novel architecture → limited tooling - Requires custom serving infrastructure - Inference 30% slower than optimized transformers initially - Less mature than transformer ecosystem

Deployment timeline:

Phase 1 (Months 1-6): Research prototype, validate components

Phase 2 (Months 7-12): Integrate and train initial model

Phase 3 (Months 13-18): Optimize for production, build tooling

Phase 4 (Months 19-24): Deploy, gather feedback, iterate

11.6 Cost-benefit analysis

Development costs:

Personnel (2 years): 3 ML researchers + 2 engineers = \$1-2M

Compute: Training + experiments = \$100-200K

Infrastructure: Serving setup = \$50K

Total: \$1.5-2.5M for full development

Potential benefits:

1. **Novel capabilities:** Explicit structure control, visual text handling 2. **Efficiency:** Hierarchical embeddings reduce memory 50-90% 3. **Multilingual:** Single model vs separate models per language 4. **Research impact:** High-novelty → top-tier publications 5. **Commercial:** Unique features for creative/design applications

ROI assessment: Positive if: - Target application requires controllable generation (creative writing, design) - Multilingual support critical (especially CJK + historical scripts) - Research prestige valued (academic/industry research lab) - Long-term platform play (not immediate commercial deployment)

12 Recommended implementation roadmap

12.1 Phase 1: Component validation (Months 1-3)

Goal: Validate each component independently

Task 1.1: Implement hierarchical embeddings

- Use Poincaré or Box embeddings on WordNet
- Benchmark memory usage and retrieval speed
- Target: 90%+ compression with <5% accuracy loss

Task 1.2: Integrate Glyph-ByT5

- Use pre-trained Glyph-ByT5-v2 for Chinese + English
- Test character-aware encoding on sample texts
- Validate 30+ point improvement on spelling tasks

Task 1.3: Train meta-structure encoders

- Grammar: Train on Universal Dependencies (100K samples)
- Meaning: Train on PropBank (100K samples)
- Emotion: Train on GoEmotions (58K samples)
- Style: Train on GYAFC + academic corpora (100K samples)
- Target: 70%+ disentanglement scores (MIG, SAP)

Task 1.4: Baseline diffusion model

- Train discrete diffusion (masked) on 1B token subset
- 512-token context, 10-50 denoising steps
- Target: Perplexity within 20% of comparable AR model

12.2 Phase 2: Integration (Months 4-9)

Goal: Combine components into unified architecture

Task 2.1: Design integrated architecture

- Define interfaces between components
- Implement 2D reshaping for diffusion
- Build conditioning mechanism (meta-structures → U-Net)

Task 2.2: End-to-end training pipeline

- Modular loss function with weight scheduling
- Gradient checkpointing for memory efficiency
- Mixed precision training (FP16 + INT8)

Task 2.3: Train initial model (1B params)

- Dataset: 50B tokens English + 25B tokens Chinese
- Hardware: 8× A100 40GB (or rent on cloud)
- Duration: ~8 weeks
- Checkpoints: Every 10B tokens for evaluation

Task 2.4: Evaluation framework

- Perplexity on held-out test sets
- Controllability: Grammar/style transfer tasks
- Multilingual: Zero-shot on Japanese, Korean
- Visual: Render text-in-image tasks

12.3 Phase 3: Optimization (Months 10-15)

Goal: Optimize for production deployment

Task 3.1: Model distillation and compression

- Distill to 500M param student model
- Quantize to INT8/INT4 (target: <2GB memory)
- Prune redundant components (30-40% sparsity)

Task 3.2: Inference optimization

- Implement efficient serving (vLLM integration)
- FlashAttention + fused kernels

- Speculative decoding (draft model)
- Target: 30-50 tokens/second on RTX 4090

Task 3.3: Batch processing and caching

- Continuous batching for multiple requests
- KV cache optimization and quantization
- Precompute meta-structure embeddings for common patterns

Task 3.4: Deployment infrastructure

- Docker containers with GPU support
- Kubernetes autoscaling
- Monitoring (Prometheus + Grafana)
- API (OpenAI-compatible endpoints)

12.4 Phase 4: Applications and iteration (Months 16-24)

Goal: Build applications, gather feedback, improve

Task 4.1: Demonstrator applications

- Creative writing assistant (style control)
- Multilingual document generation
- Historical text decipherment tool
- Typography/design generation

Task 4.2: User studies and feedback

- 100-1000 users testing applications
- Collect preference data (ILM vs GPT-4)
- Identify failure modes and limitations

Task 4.3: Targeted improvements

- Fine-tune on domains where ILM underperforms
- Add new meta-structures based on user needs
- Extend to additional languages/scripts

Task 4.4: Research publications

- Architecture paper (ICLR/NeurIPS)
- Multilingual visual encoding (ACL/EMNLP)
- Applications and case studies (domain conferences)

12.5 Minimum viable implementation

For proof-of-concept (6-month timeline, 2 researchers):

Scope reduction:

- Skip hierarchical embeddings initially (use standard embeddings)
- Use pre-trained Glyph-ByT5 (don't train from scratch)
- Simpler meta-structures (grammar + sentiment only, 50 patterns total)
- Smaller model (300M params)
- Limited languages (English + Chinese only)

- Shorter context (256 tokens vs 512-1024)

Resources:

- Compute: 4× A100 40GB for 4 weeks (cloud: ~\$10-15K)
- Data: Use existing datasets (no custom annotation)
- Personnel: 2 researchers + 1 engineer

Deliverable:

- Working prototype demonstrating core concept
- Benchmark against standard diffusion LM
- Demonstrate controllable generation
- Sufficient for research paper submission

Key insights Superposition enables control; hierarchical structure enables compression and interpretability; diffusion enables global revision; and visual encoding enables universal multilingual handling. Together, ILM offers controllable, efficient, and extensible language generation.

Appendix: Mathematical Notes

A. Graph Laplacian

With affinity W and degree matrix Deg , $L = \text{Deg} - W$ yields $\text{tr}(E^\top L E) = \frac{1}{2} \sum_{p,q} w_{pq} \|E_p - E_q\|_2^2$.

B. Mixed-Radix Posterior Capacity

Given $C = \prod K_\ell$ codes and V items, balanced assignment reduces expected prefix collisions; nearest-neighbor lookup on prefixes accelerates decoding and improves robustness to denoise errors.